

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.formula.api as smf
import statsmodels.api as sm
from statsmodels.graphics.regressionplots import influence_plot
```

```
In [2]: data = pd.read_csv('ToyotaCorolla.csv', encoding = 'latin1')
data.head()
```

Out[2]:

	Id	Model	Price	Age_08_04	Mfg_Month	Mfg_Year	KM	Fuel_Type	HP	Met_Colo
0	1	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3- Doors	13500	23	10	2002	46986	Diesel	90	.
1	2	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3- Doors	13750	23	10	2002	72937	Diesel	90	.
2	3	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3- Doors	13950	24	9	2002	41711	Diesel	90	.
3	4	TOYOTA Corolla 2.0 D4D HATCHB TERRA 2/3- Doors	14950	26	7	2002	48000	Diesel	90	(
4	5	TOYOTA Corolla 2.0 D4D HATCHB SOL 2/3- Doors	13750	30	3	2002	38500	Diesel	90	(

5 rows × 38 columns

```
In [3]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1436 entries, 0 to 1435
Data columns (total 38 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Id                                     1436 non-null   int64
1   Model                                1436 non-null   object
2   Price                                1436 non-null   int64
3   Age_08_04                            1436 non-null   int64
4   Mfg_Month                            1436 non-null   int64
5   Mfg_Year                             1436 non-null   int64
6   KM                                    1436 non-null   int64
7   Fuel_Type                            1436 non-null   object
8   HP                                    1436 non-null   int64
9   Met_Color                            1436 non-null   int64
10  Color                                1436 non-null   object
11  Automatic                            1436 non-null   int64
12  cc                                    1436 non-null   int64
13  Doors                                1436 non-null   int64
14  Cylinders                            1436 non-null   int64
15  Gears                                1436 non-null   int64
16  Quarterly_Tax                        1436 non-null   int64
17  Weight                                1436 non-null   int64
18  Mfr_Guarantee                        1436 non-null   int64
19  BOVAG_Guarantee                      1436 non-null   int64
20  Guarantee_Period                    1436 non-null   int64
21  ABS                                  1436 non-null   int64
22  Airbag_1                             1436 non-null   int64
23  Airbag_2                             1436 non-null   int64
24  Airco                                1436 non-null   int64
25  Automatic_airco                      1436 non-null   int64
26  Boardcomputer                        1436 non-null   int64
27  CD_Player                            1436 non-null   int64
28  Central_Lock                         1436 non-null   int64
29  Powered_Windows                      1436 non-null   int64
30  Power_Steering                       1436 non-null   int64
31  Radio                                1436 non-null   int64
32  Mistlamps                            1436 non-null   int64
33  Sport_Model                          1436 non-null   int64
34  Backseat_Divider                     1436 non-null   int64
35  Metallic_Rim                        1436 non-null   int64
36  Radio_cassette                       1436 non-null   int64
37  Tow_Bar                              1436 non-null   int64
dtypes: int64(35), object(3)
memory usage: 426.4+ KB
```

In [4]: `data.describe()`

Out [4]:

	Id	Price	Age_08_04	Mfg_Month	Mfg_Year	KM
<b>count</b>	1436.000000	1436.000000	1436.000000	1436.000000	1436.000000	1436.000000
<b>mean</b>	721.555014	10730.824513	55.947075	5.548747	1999.625348	68533.259749
<b>std</b>	416.476890	3626.964585	18.599988	3.354085	1.540722	37506.448872
<b>min</b>	1.000000	4350.000000	1.000000	1.000000	1998.000000	1.000000
<b>25%</b>	361.750000	8450.000000	44.000000	3.000000	1998.000000	43000.000000
<b>50%</b>	721.500000	9900.000000	61.000000	5.000000	1999.000000	63389.500000
<b>75%</b>	1081.250000	11950.000000	70.000000	8.000000	2001.000000	87020.750000
<b>max</b>	1442.000000	32500.000000	80.000000	12.000000	2004.000000	243000.000000

8 rows × 35 columns

In [5]: `data1=data.iloc[:, [2,3,6,8,12,13,15,16,17]]`  
`data1.dtypes`

Out [5]: Price int64  
 Age\_08\_04 int64  
 KM int64  
 HP int64  
 cc int64  
 Doors int64  
 Gears int64  
 Quarterly\_Tax int64  
 Weight int64  
 dtype: object

In [6]: `data2=data1.rename({"Age_08_04":"age", "Quarterly_Tax":"tax", "Weight"`  
`data2.head()`

Out [6]:

	Price	age	KM	HP	cc	door	Gears	tax	weight
<b>0</b>	13500	23	46986	90	2000	3	5	210	1165
<b>1</b>	13750	23	72937	90	2000	3	5	210	1165
<b>2</b>	13950	24	41711	90	2000	3	5	210	1165
<b>3</b>	14950	26	48000	90	2000	3	5	210	1165
<b>4</b>	13750	30	38500	90	2000	3	5	210	1170

```
In [7]: data2.isnull().sum()
```

```
Out[7]: Price      0
age      0
KM      0
HP      0
cc      0
door     0
Gears    0
tax      0
weight   0
dtype: int64
```

```
In [8]: data2.duplicated().value_counts()
data2[data2.duplicated()]
```

```
Out[8]:
```

	Price	age	KM	HP	cc	door	Gears	tax	weight
113	24950	8	13253	116	2000	5	5	234	1320

```
In [9]: data2.drop_duplicates(inplace=True)
```


```
In [10]: data2.corr()
```

```
Out[10]:
```

	Price	age	KM	HP	cc	door	Gears	tax	weight
<b>Price</b>	1.000000	-0.876273	-0.569420	0.314134	0.124375	0.183604	0.063831	0.211508	0.575869
<b>age</b>	-0.876273	1.000000	0.504575	-0.155293	-0.096549	-0.146929	-0.005629	-0.193319	-0.466484
<b>KM</b>	-0.569420	0.504575	1.000000	-0.332904	0.103822	-0.035193	0.014890	0.283312	-0.023969
<b>HP</b>	0.314134	-0.155293	-0.332904	1.000000	0.035207	0.091803	0.209642	-0.302287	0.087143
<b>cc</b>	0.124375	-0.096549	0.103822	0.035207	1.000000	0.079254	0.014732	0.305982	0.335077
<b>door</b>	0.183604	-0.146929	-0.035193	0.091803	0.079254	1.000000	-0.160101	0.107353	0.301734
<b>Gears</b>	0.063831	-0.005629	0.014890	0.209642	0.014732	-0.160101	1.000000	-0.005125	0.021238
<b>tax</b>	0.211508	-0.193319	0.283312	-0.302287	0.305982	0.107353	-0.005125	1.000000	0.621961
<b>weight</b>	0.575869	-0.466484	-0.023969	0.087143	0.335077	0.301734	0.021238	0.621961	1.000000

```
In [11]: sns.set_style(style="darkgrid")  
sns.pairplot(data= data2)
```

```
Out[11]: <seaborn.axisgrid.PairGrid at 0x7ffd7683ed00>
```



removing gears column as it has very less correlation with Price

```
In [12]: data3 = data2.drop("Gears", axis=1)
data3.head()
```

Out[12]:

	Price	age	KM	HP	cc	door	tax	weight
0	13500	23	46986	90	2000	3	210	1165
1	13750	23	72937	90	2000	3	210	1165
2	13950	24	41711	90	2000	3	210	1165
3	14950	26	48000	90	2000	3	210	1165
4	13750	30	38500	90	2000	3	210	1170

```
In [13]: data3.corr()
```

Out[13]:

	Price	age	KM	HP	cc	door	tax	weig
Price	1.000000	-0.876273	-0.569420	0.314134	0.124375	0.183604	0.211508	0.57589
age	-0.876273	1.000000	0.504575	-0.155293	-0.096549	-0.146929	-0.193319	-0.46648
KM	-0.569420	0.504575	1.000000	-0.332904	0.103822	-0.035193	0.283312	-0.02396
HP	0.314134	-0.155293	-0.332904	1.000000	0.035207	0.091803	-0.302287	0.08714
cc	0.124375	-0.096549	0.103822	0.035207	1.000000	0.079254	0.305982	0.33507
door	0.183604	-0.146929	-0.035193	0.091803	0.079254	1.000000	0.107353	0.30173
tax	0.211508	-0.193319	0.283312	-0.302287	0.305982	0.107353	1.000000	0.62198
weight	0.575869	-0.466484	-0.023969	0.087143	0.335077	0.301734	0.621988	1.00000

```
In [14]: model=smf.ols("Price~age+KM+HP+cc+door+tax+weight",data = data3).fit()
model.summary()
```

Out [14]: OLS Regression Results

<b>Dep. Variable:</b>	Price	<b>R-squared:</b>	0.862
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.861
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	1269.
<b>Date:</b>	Wed, 07 Dec 2022	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	01:16:36	<b>Log-Likelihood:</b>	-12371.
<b>No. Observations:</b>	1435	<b>AIC:</b>	2.476e+04
<b>Df Residuals:</b>	1427	<b>BIC:</b>	2.480e+04
<b>Df Model:</b>	7		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	-2636.3843	1061.677	-2.483	0.013	-4718.999	-553.770
<b>age</b>	-121.8478	2.622	-46.476	0.000	-126.991	-116.705
<b>KM</b>	-0.0205	0.001	-16.325	0.000	-0.023	-0.018
<b>HP</b>	33.6479	2.742	12.270	0.000	28.269	39.027
<b>cc</b>	-0.1227	0.090	-1.358	0.175	-0.300	0.054
<b>door</b>	-23.9806	39.372	-0.609	0.543	-101.214	53.253
<b>tax</b>	4.0353	1.313	3.073	0.002	1.459	6.611
<b>weight</b>	16.8844	1.072	15.743	0.000	14.781	18.988

<b>Omnibus:</b>	149.646	<b>Durbin-Watson:</b>	1.551
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	980.896
<b>Skew:</b>	-0.217	<b>Prob(JB):</b>	1.00e-213
<b>Kurtosis:</b>	7.027	<b>Cond. No.</b>	2.34e+06

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.34e+06. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [15]: m1 = smf.ols('Price~cc', data=data3).fit()
m1.summary()
```

Out [15]: OLS Regression Results

<b>Dep. Variable:</b>	Price	<b>R-squared:</b>	0.015
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.015
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	22.52
<b>Date:</b>	Wed, 07 Dec 2022	<b>Prob (F-statistic):</b>	2.29e-06
<b>Time:</b>	01:16:36	<b>Log-Likelihood:</b>	-13779.
<b>No. Observations:</b>	1435	<b>AIC:</b>	2.756e+04
<b>Df Residuals:</b>	1433	<b>BIC:</b>	2.757e+04
<b>Df Model:</b>	1		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	9053.5368	363.894	24.880	0.000	8339.715	9767.359
<b>cc</b>	1.0576	0.223	4.745	0.000	0.620	1.495

<b>Omnibus:</b>	463.846	<b>Durbin-Watson:</b>	0.269
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	1386.822
<b>Skew:</b>	1.645	<b>Prob(JB):</b>	7.17e-302
<b>Kurtosis:</b>	6.518	<b>Cond. No.</b>	6.28e+03

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 6.28e+03. This might indicate that there are strong multicollinearity or other numerical problems.



```
In [16]: m1 = smf.ols('Price~door', data= data3).fit()
m1.summary()
```

Out[16]: OLS Regression Results

<b>Dep. Variable:</b>	Price	<b>R-squared:</b>	0.034
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.033
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	49.99
<b>Date:</b>	Wed, 07 Dec 2022	<b>Prob (F-statistic):</b>	2.40e-12
<b>Time:</b>	01:16:36	<b>Log-Likelihood:</b>	-13765.
<b>No. Observations:</b>	1435	<b>AIC:</b>	2.753e+04
<b>Df Residuals:</b>	1433	<b>BIC:</b>	2.755e+04
<b>Df Model:</b>	1		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	7916.1452	407.596	19.422	0.000	7116.596	8715.694
<b>door</b>	695.4978	98.366	7.071	0.000	502.541	888.454

<b>Omnibus:</b>	465.543	<b>Durbin-Watson:</b>	0.289
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	1403.980
<b>Skew:</b>	1.647	<b>Prob(JB):</b>	1.35e-305
<b>Kurtosis:</b>	6.554	<b>Cond. No.</b>	19.0

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [17]: data4 = data3.drop("cc", axis = 1)
data4.head()
```

Out[17]:

	Price	age	KM	HP	door	tax	weight
<b>0</b>	13500	23	46986	90	3	210	1165
<b>1</b>	13750	23	72937	90	3	210	1165
<b>2</b>	13950	24	41711	90	3	210	1165
<b>3</b>	14950	26	48000	90	3	210	1165
<b>4</b>	13750	30	38500	90	3	210	1170

In [18]: `data4.corr()`

Out[18]:

	Price	age	KM	HP	door	tax	weight
Price	1.000000	-0.876273	-0.569420	0.314134	0.183604	0.211508	0.575869
age	-0.876273	1.000000	0.504575	-0.155293	-0.146929	-0.193319	-0.466484
KM	-0.569420	0.504575	1.000000	-0.332904	-0.035193	0.283312	-0.023969
HP	0.314134	-0.155293	-0.332904	1.000000	0.091803	-0.302287	0.087143
door	0.183604	-0.146929	-0.035193	0.091803	1.000000	0.107353	0.301734
tax	0.211508	-0.193319	0.283312	-0.302287	0.107353	1.000000	0.621988
weight	0.575869	-0.466484	-0.023969	0.087143	0.301734	0.621988	1.000000

```
In [19]: m3= smf.ols('Price~age+KM+HP+door+tax+weight',data=data4).fit()
m3.summary()
```

Out [19]: OLS Regression Results

<b>Dep. Variable:</b>	Price	<b>R-squared:</b>	0.861
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.861
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	1480.
<b>Date:</b>	Wed, 07 Dec 2022	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	01:16:37	<b>Log-Likelihood:</b>	-12372.
<b>No. Observations:</b>	1435	<b>AIC:</b>	2.476e+04
<b>Df Residuals:</b>	1428	<b>BIC:</b>	2.479e+04
<b>Df Model:</b>	6		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	-2518.7574	1058.453	-2.380	0.017	-4595.046	-442.468
<b>age</b>	-121.8917	2.622	-46.483	0.000	-127.036	-116.748
<b>KM</b>	-0.0206	0.001	-16.455	0.000	-0.023	-0.018
<b>HP</b>	33.3055	2.731	12.194	0.000	27.947	38.663
<b>door</b>	-23.2523	39.380	-0.590	0.555	-100.502	53.997
<b>tax</b>	3.8184	1.304	2.929	0.003	1.261	6.376
<b>weight</b>	16.6515	1.059	15.724	0.000	14.574	18.729

<b>Omnibus:</b>	145.933	<b>Durbin-Watson:</b>	1.548
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	937.411
<b>Skew:</b>	-0.206	<b>Prob(JB):</b>	2.78e-204
<b>Kurtosis:</b>	6.938	<b>Cond. No.</b>	2.33e+06

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.33e+06. This might indicate that there are strong multicollinearity or other numerical problems.

## vif calculation

```
In [20]: rsq_age=smf.ols('age~KM+HP+door+tax+weight',data=data4).fit().rsqua
vif_age=1/(1-rsq_age)
```

```
In [21]: rsq_KM=smf.ols('KM~age+HP+door+tax+weight',data=data4).fit().rsqua
vif_KM=1/(1-rsq_KM)
```

```
In [22]: rsq_HP=smf.ols('HP~age+KM+door+tax+weight',data=data4).fit().rsqua
vif_HP=1/(1-rsq_HP)
```

```
In [23]: rsq_door=smf.ols('door~age+KM+HP+tax+weight',data=data4).fit().rsqu
vif_door=1/(1-rsq_door)
```

```
In [24]: rsq_tax=smf.ols('tax~age+KM+HP+door+weight',data=data4).fit().rsqua
vif_tax=1/(1-rsq_tax)
```

```
In [25]: rsq_Weight=smf.ols('weight~age+KM+HP+door+tax',data=data4).fit().rs
vif_Weight=1/(1-rsq_Weight)
```

```
In [26]: d1 = {'variables':['age','KM','HP','door','tax','weight'],
               'vif':[vif_age,vif_KM,vif_HP,vif_door,vif_tax,vif_Weight]}
```

```
In [27]: vif_frame = pd.DataFrame(d1)
vif_frame
```

Out[27]:

	variables	vif
0	age	1.875416
1	KM	1.738917
2	HP	1.325246
3	door	1.113939
4	tax	2.256997
5	weight	2.423423

## residual analysis for normality

```
In [28]: res=m3.resid  
res.head(10)
```

```
Out[28]: 0    -3339.719690  
1    -2555.830091  
2    -2876.350450  
3    -1503.183434  
4    -2494.317424  
5    -2587.641707  
6     195.534750  
7    1876.021951  
8    1276.335218  
9    -1598.873899  
dtype: float64
```


```
In [29]: res.mean()
```

```
Out[29]: 1.042505291384687e-09
```

```
In [30]: import statsmodels.api as sm  
qqplot=sm.qqplot(res,line='q')  
plt.title('test for normality of residuals(q-qplot)')  
plt.show
```

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/graphics/gofplots.py:993: UserWarning: marker is redundantly defined by the 'marker' keyword argument and the fmt string "bo" (-> marker='o'). The keyword argument will take precedence.  
ax.plot(x, y, fmt, \*\*plot\_style)


```
Out[30]: <function matplotlib.pyplot.show(close=None, block=None)>
```



```
In [31]: sns.boxplot(m3.resid)
```

```
/opt/anaconda3/lib/python3.9/site-packages/seaborn/_decorators.py:
36: FutureWarning: Pass the following variable as a keyword arg: x
. From version 0.12, the only valid positional argument will be `d
ata`, and passing other arguments without an explicit keyword will
result in an error or misinterpretation.
    warnings.warn(
```

```
Out [31]: <AxesSubplot:>
```



```
In [32]: list(np.where(m3.resid < -6000))
```


```
Out [32]: [array([220, 600, 959])]
```

## residual plot for homoscedasticity

```
In [33]: def get_standardized_values(vals):
          return (vals - vals.mean()) / vals.std()
```

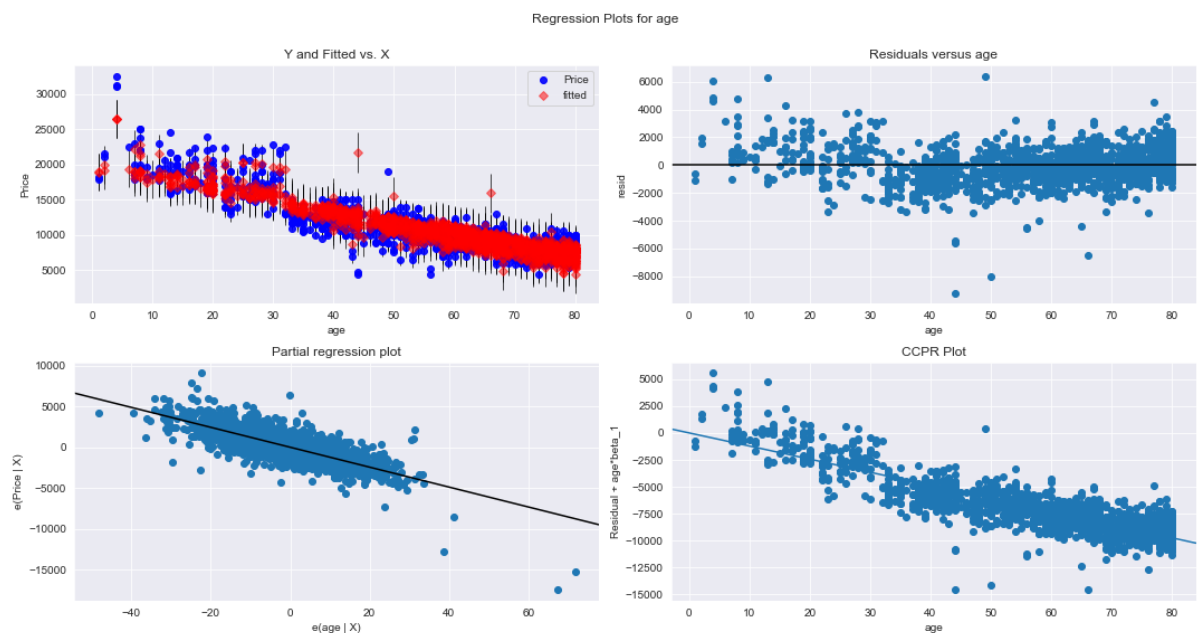
```
In [34]: plt.scatter(get_standardized_values(m3.fittedvalues),
                     get_standardized_values(m3.resid))
```

```
Out[34]: <matplotlib.collections.PathCollection at 0x7ffd79bd3490>
```



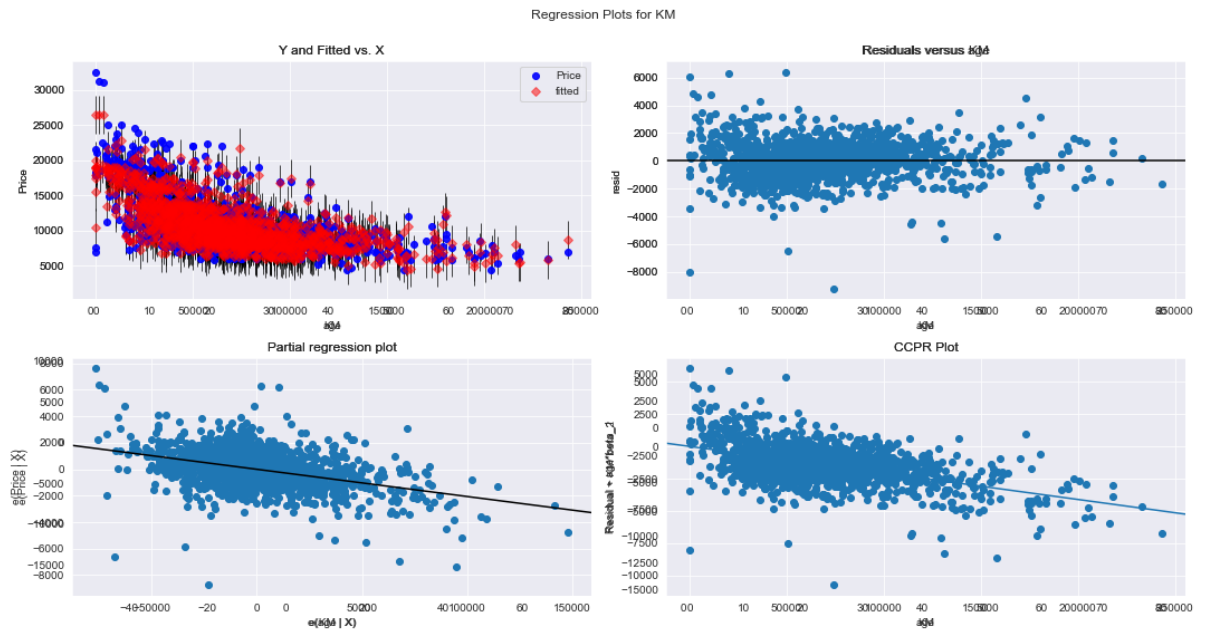
## residual vs regressor

```
In [35]: fig = plt.figure(figsize=(15,8))
fig = sm.graphics.plot_regress_exog(m3, 'age', fig=fig)
plt.show()
```

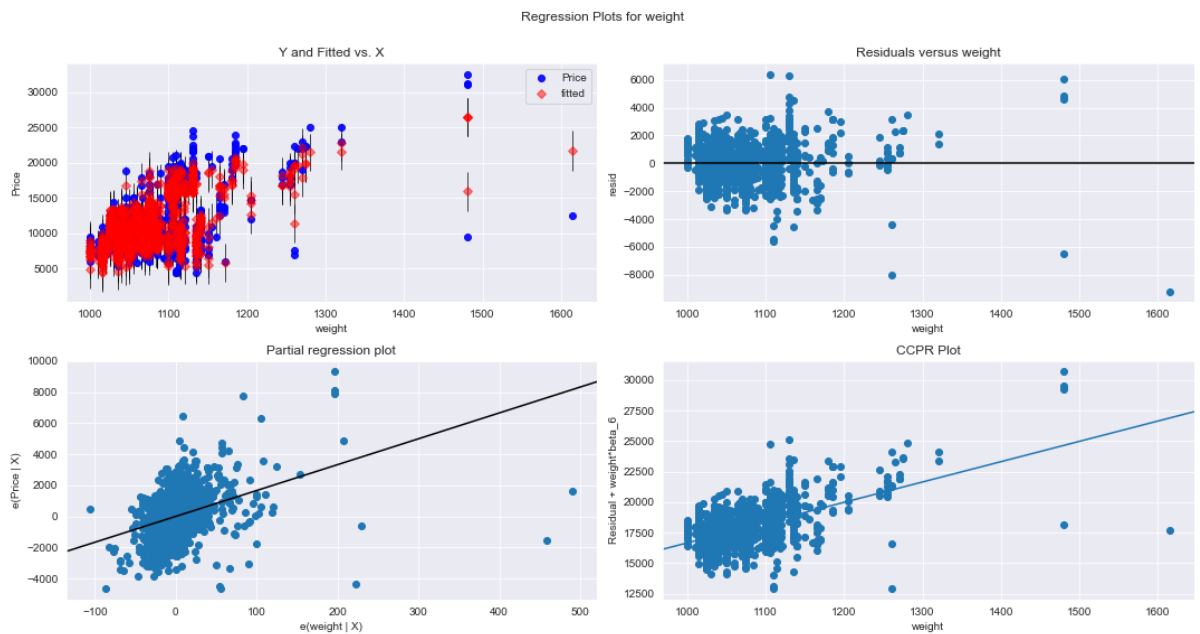


```
In [36]: fig = sm.graphics.plot_regress_exog(m3, 'KM', fig=fig)
fig
```

Out [36]:



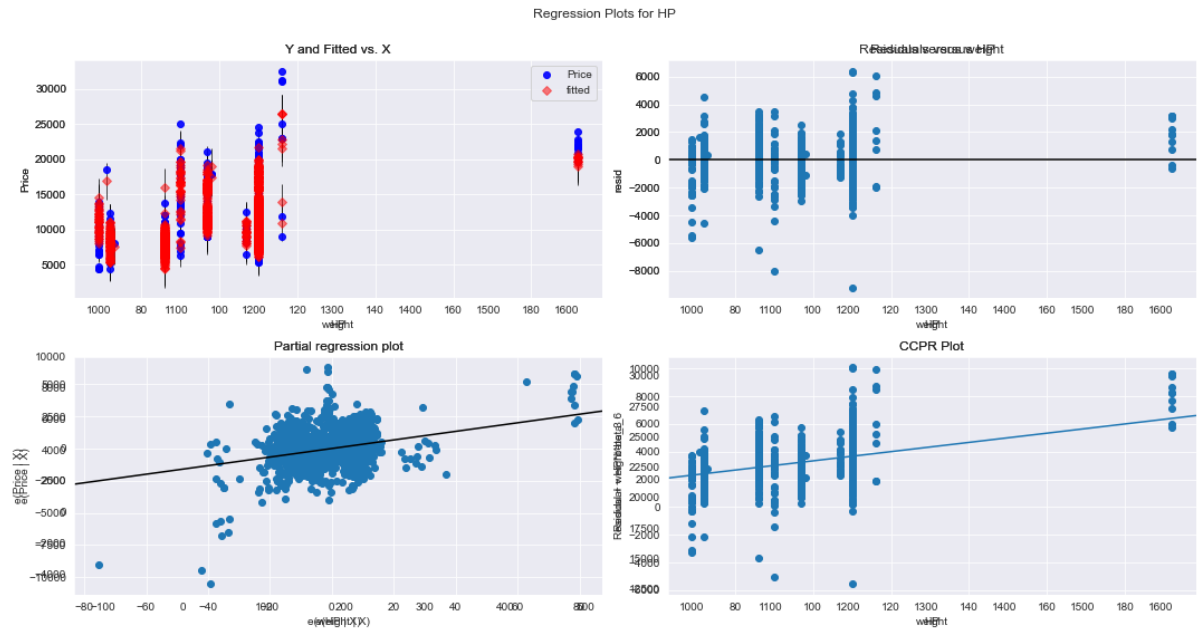
```
In [37]: fig = plt.figure(figsize=(15,8))
fig = sm.graphics.plot_regress_exog(m3, 'weight', fig=fig)
plt.show()
```





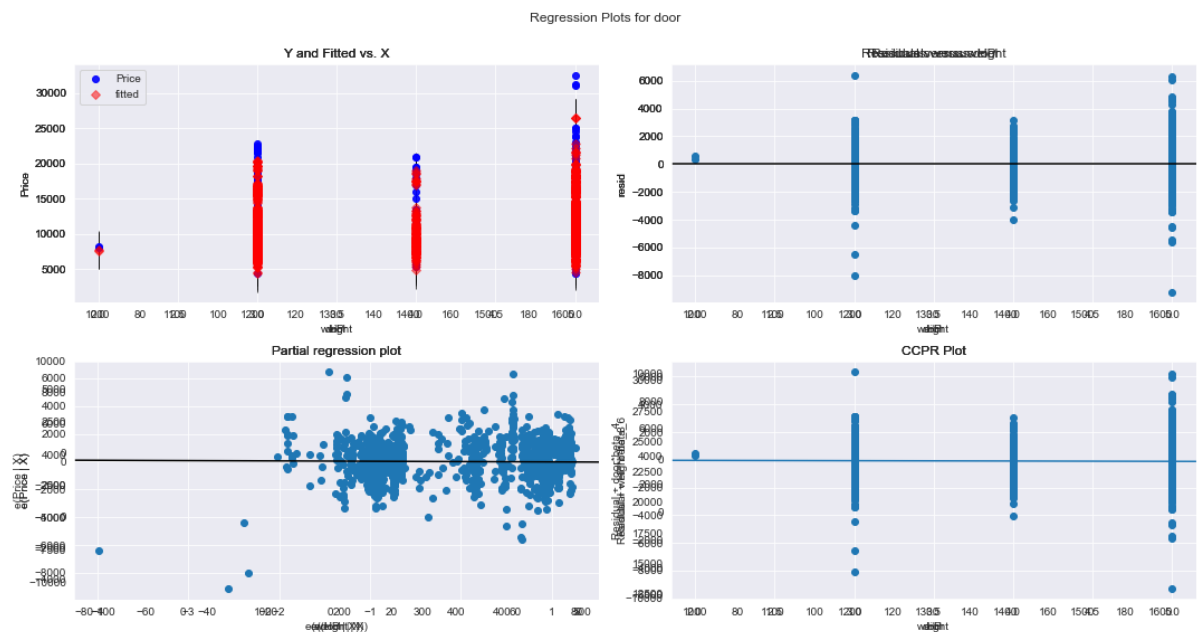
In [38]: `sm.graphics.plot_regress_exog(m3, 'HP', fig=fig)`

Out [38]:



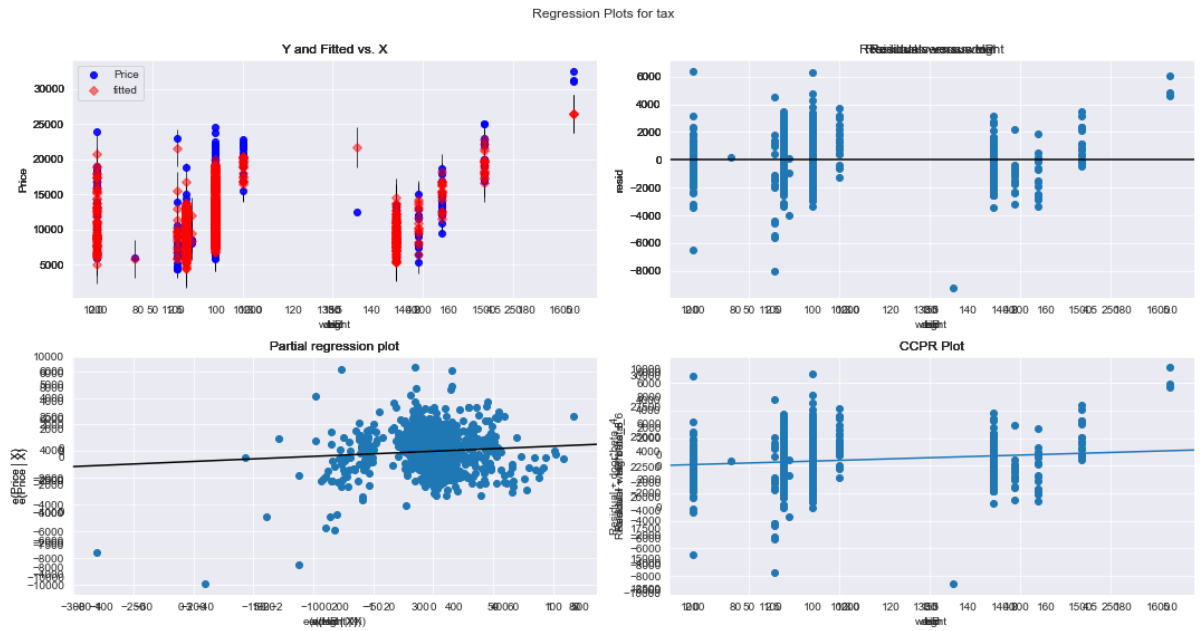
In [39]: `sm.graphics.plot_regress_exog(m3, 'door', fig=fig)`

Out [39]:



```
In [40]: sm.graphics.plot_regress_exog(m3, 'tax', fig=fig)
```

Out[40]:




```
In [41]: model_influence=m3.get_influence()
(c, _) = model_influence.cooks_distance
```

In [42]: c

Out[42]: array([9.38906054e-03, 5.17943805e-03, 7.10668449e-03, ...,  
2.10208718e-06, 8.65451174e-04, 1.37001373e-02])

```
In [43]: fig=plt.subplots(figsize=(20,7))
plt.stem(np.arange(len(data4)),np.round(c,3))
plt.xlabel("index")
plt.ylabel("cooks distance")
```

Out[43]: Text(0, 0.5, 'cooks distance')



```
In [44]: np.argmax(c),np.max(c)
```

```
Out[44]: (220, 1.1987544668622254)
```

```
In [45]: data_new=data4.drop(data4.index[220],axis=0)  
data_new.head()
```

```
Out[45]:
```

	Price	age	KM	HP	door	tax	weight
0	13500	23	46986	90	3	210	1165
1	13750	23	72937	90	3	210	1165
2	13950	24	41711	90	3	210	1165
3	14950	26	48000	90	3	210	1165
4	13750	30	38500	90	3	210	1170

```
In [46]: m4 = smf.ols("Price~age+KM+HP+door+tax+weight",data=data_new).fit()
m4.summary()
```

Out [46]: OLS Regression Results

<b>Dep. Variable:</b>	Price	<b>R-squared:</b>	0.867
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.866
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	1546.
<b>Date:</b>	Wed, 07 Dec 2022	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	01:16:50	<b>Log-Likelihood:</b>	-12336.
<b>No. Observations:</b>	1434	<b>AIC:</b>	2.469e+04
<b>Df Residuals:</b>	1427	<b>BIC:</b>	2.472e+04
<b>Df Model:</b>	6		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	-5450.5835	1109.915	-4.911	0.000	-7627.823	-3273.344
<b>age</b>	-119.1731	2.599	-45.861	0.000	-124.271	-114.076
<b>KM</b>	-0.0208	0.001	-16.937	0.000	-0.023	-0.018
<b>HP</b>	31.5845	2.690	11.741	0.000	26.308	36.861
<b>door</b>	-46.5493	38.767	-1.201	0.230	-122.595	29.497
<b>tax</b>	1.9212	1.304	1.473	0.141	-0.637	4.479
<b>weight</b>	19.6687	1.114	17.648	0.000	17.482	21.855


<b>Omnibus:</b>	131.540	<b>Durbin-Watson:</b>	1.580
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	729.534
<b>Skew:</b>	-0.206	<b>Prob(JB):</b>	3.83e-159
<b>Kurtosis:</b>	6.470	<b>Cond. No.</b>	2.49e+06

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.49e+06. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [51]: model_influence_1 = m4.get_influence()
(c1, _) = model_influence_1.cooks_distance
```

```
In [52]: fig = plt.subplots(figsize = (20,7))
plt.stem(np.arange(len(data_new)), np.round(c1,3))
plt.xlabel("Row Index")
plt.ylabel("Cooks Distance")
plt.show()
```



```
In [54]: np.argmax(c1), np.max(c1)
```

```
Out[54]: (958, 1.3402996914231093)
```

```
In [55]: dataa = data_new.drop(data_new.index[958],axis=0).reset_index()
```

```
In [56]: dataa
```

```
Out[56]:
```

	index	Price	age	KM	HP	door	tax	weight
0	0	13500	23	46986	90	3	210	1165
1	1	13750	23	72937	90	3	210	1165
2	2	13950	24	41711	90	3	210	1165
3	3	14950	26	48000	90	3	210	1165
4	4	13750	30	38500	90	3	210	1170
...	...	...	...	...	...	...	...	...
1428	1431	7500	69	20544	86	3	69	1025
1429	1432	10845	72	19000	86	3	69	1015
1430	1433	8500	71	17016	86	3	69	1015
1431	1434	7250	70	16916	86	3	69	1015
1432	1435	6950	76	1	110	5	19	1114

1433 rows × 8 columns

```
In [57]: dataa1 = data_new.drop(data_new.index[958],axis=0).reset_index()
```

```
In [59]: dataa1.drop(['index'],axis = 1)
```

Out[59]:

	Price	age	KM	HP	door	tax	weight
0	13500	23	46986	90	3	210	1165
1	13750	23	72937	90	3	210	1165
2	13950	24	41711	90	3	210	1165
3	14950	26	48000	90	3	210	1165
4	13750	30	38500	90	3	210	1170
...	...	...	...	...	...	...	...
1428	7500	69	20544	86	3	69	1025
1429	10845	72	19000	86	3	69	1015
1430	8500	71	17016	86	3	69	1015
1431	7250	70	16916	86	3	69	1015
1432	6950	76	1	110	5	19	1114

1433 rows × 7 columns

```
In [60]: m5 = smf.ols("Price~age+KM+HP+door+tax+weight",data=dataa1).fit()
m5.summary()
```

Out [60]: OLS Regression Results

<b>Dep. Variable:</b>	Price	<b>R-squared:</b>	0.871
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.870
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	1602.
<b>Date:</b>	Wed, 07 Dec 2022	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	01:36:40	<b>Log-Likelihood:</b>	-12305.
<b>No. Observations:</b>	1433	<b>AIC:</b>	2.462e+04
<b>Df Residuals:</b>	1426	<b>BIC:</b>	2.466e+04
<b>Df Model:</b>	6		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	-8477.4104	1181.676	-7.174	0.000	-1.08e+04	-6159.400
<b>age</b>	-116.1584	2.598	-44.713	0.000	-121.254	-111.062
<b>KM</b>	-0.0211	0.001	-17.446	0.000	-0.023	-0.019
<b>HP</b>	28.3392	2.692	10.525	0.000	23.058	33.621
<b>door</b>	-82.8874	38.556	-2.150	0.032	-158.519	-7.255
<b>tax</b>	-0.8966	1.351	-0.664	0.507	-3.546	1.753
<b>weight</b>	23.0339	1.206	19.103	0.000	20.669	25.399


<b>Omnibus:</b>	135.138	<b>Durbin-Watson:</b>	1.631
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	741.678
<b>Skew:</b>	-0.235	<b>Prob(JB):</b>	8.84e-162
<b>Kurtosis:</b>	6.493	<b>Cond. No.</b>	2.69e+06

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.69e+06. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [61]: model_influence_2 = m5.get_influence()
(c2, _) = model_influence_2.cooks_distance
```

```
In [63]: fig = plt.subplots(figsize = (20,7))
plt.stem(np.arange(len(dataa1)), np.round(c2,3))
plt.xlabel("Row Index")
plt.ylabel("Cooks Distance")
plt.show()
```



```
In [64]: np.argmax(c2), np.max(c2)
```

```
Out[64]: (599, 0.4074330373473679)
```

```
In [65]: data.shape
```

```
Out[65]: (1436, 38)
```

```
In [66]: k = data.shape[0]
n = data.shape[1]
lev_cutoff = 3*((k+1)/n)
lev_cutoff
```


```
Out[66]: 113.44736842105263
```

```
In [67]: from statsmodels.graphics.regressionplots import influence_plot
```



```
In [73]: influence_plot(model, alpha = 0.5)

y= [i for i in range (-2,8)]
x= [lev_cutoff for i in range(10)]
plt.plot(x,y, 'r+')
plt.show()
```



```
In [74]: final_model = smf.ols("Price~age+KM+HP+door+tax+weight",data=dataa1)
```


```
In [75]: final_model.rsquared , final_model.aic
```

```
Out[75]: (0.8707773883293434, 24624.807746446797)
```

```
In [76]: pred_y = final_model.predict(dataa1)
pred_y
```

```
Out[76]: 0      16807.692483
1      16260.176738
2      16802.826307
3      16437.823729
4      16288.791032
...
1428     8810.603810
1429     8264.364896
1430     8422.381886
1431     8540.650133
1432    11040.126985
Length: 1433, dtype: float64
```

```
In [77]: x1 = dataa1  
y1 = pred_y  
plt.title('Actual Values vs Predicted values')  
plt.xlabel("Predicted Values")  
plt.ylabel(" Actual Values")  
plt.plot(x1,y1)  
plt.show()
```



```
In [ ]:
```