

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import statsmodels.formula.api as smf
import matplotlib.pyplot as plt
from sklearn import linear_model
```

```
In [2]: df=pd.read_csv('delivery_time.csv')
df
```

Out[2]:

	Delivery Time	Sorting Time
0	21.00	10
1	13.50	4
2	19.75	6
3	24.00	9
4	29.00	10
5	15.35	6
6	19.00	7
7	9.50	3
8	17.90	10
9	18.75	9
10	19.83	8
11	10.75	4
12	16.68	7
13	11.50	3
14	12.03	3
15	14.88	4
16	13.75	6
17	18.11	7
18	8.00	2
19	17.83	7
20	21.50	5

In [3]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21 entries, 0 to 20
Data columns (total 2 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   Delivery Time    21 non-null     float64
 1   Sorting Time     21 non-null     int64   
dtypes: float64(1), int64(1)
memory usage: 464.0 bytes
```

In [4]: `df.describe()`

Out [4]:

	Delivery Time	Sorting Time
count	21.000000	21.000000
mean	16.790952	6.190476
std	5.074901	2.542028
min	8.000000	2.000000
25%	13.500000	4.000000
50%	17.830000	6.000000
75%	19.750000	8.000000
max	29.000000	10.000000

In [38]: `df.dtypes`

Out [38]:

Delivery Time	float64
Sorting Time	int64
Deliverd_predict	float64
dtype:	object

```
In [39]: df.isnull()
```

```
Out[39]:
```

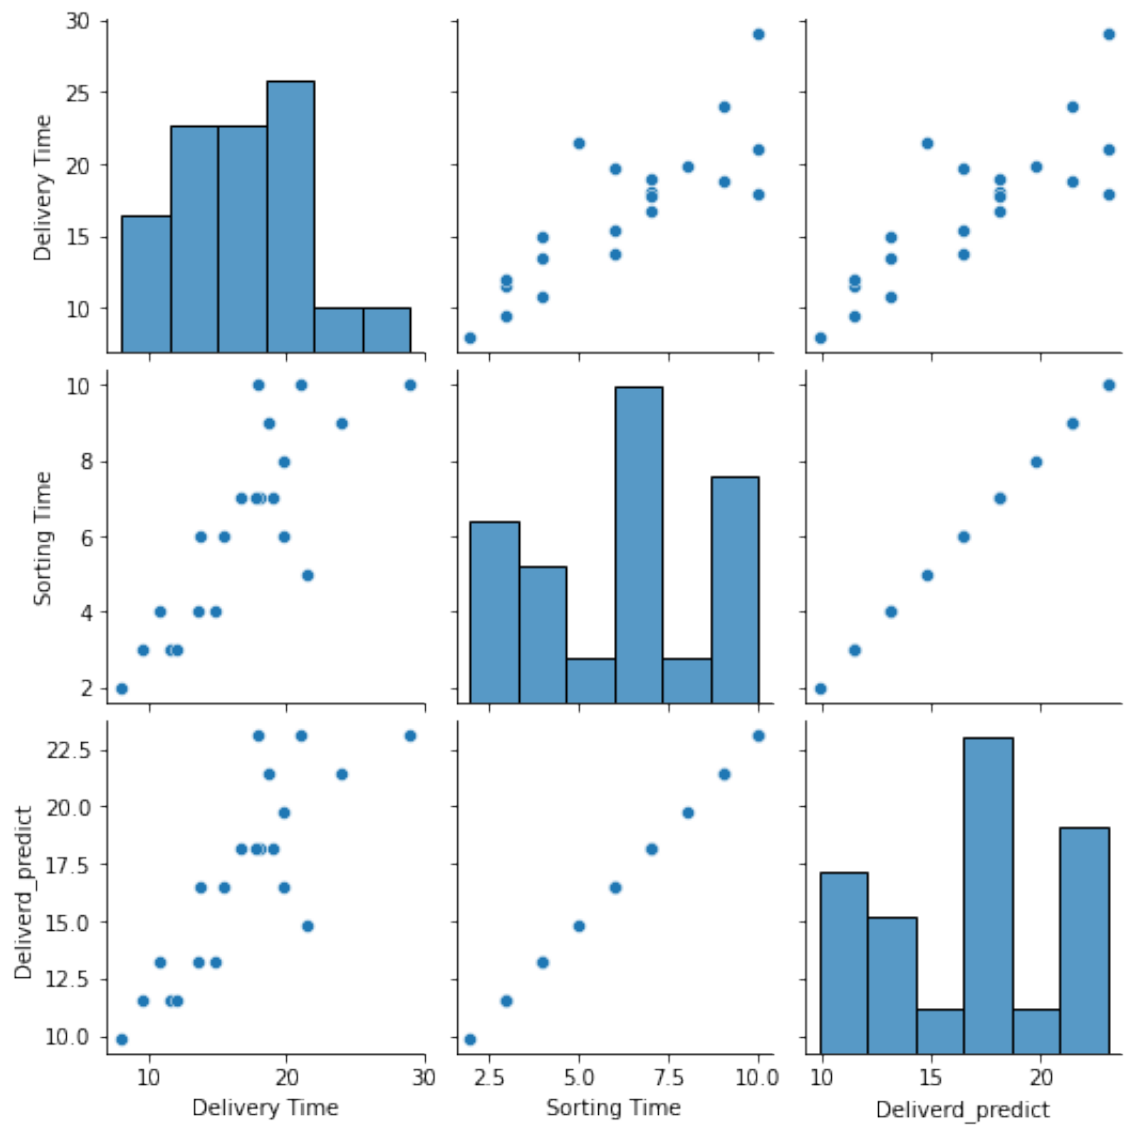
	Delivery Time	Sorting Time	Deliverd_predict
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	False
5	False	False	False
6	False	False	False
7	False	False	False
8	False	False	False
9	False	False	False
10	False	False	False
11	False	False	False
12	False	False	False
13	False	False	False
14	False	False	False
15	False	False	False
16	False	False	False
17	False	False	False
18	False	False	False
19	False	False	False
20	False	False	False

```
In [41]: df.isnull().sum()
```

```
Out[41]: Delivery Time      0
Sorting Time      0
Deliverd_predict  0
dtype: int64
```

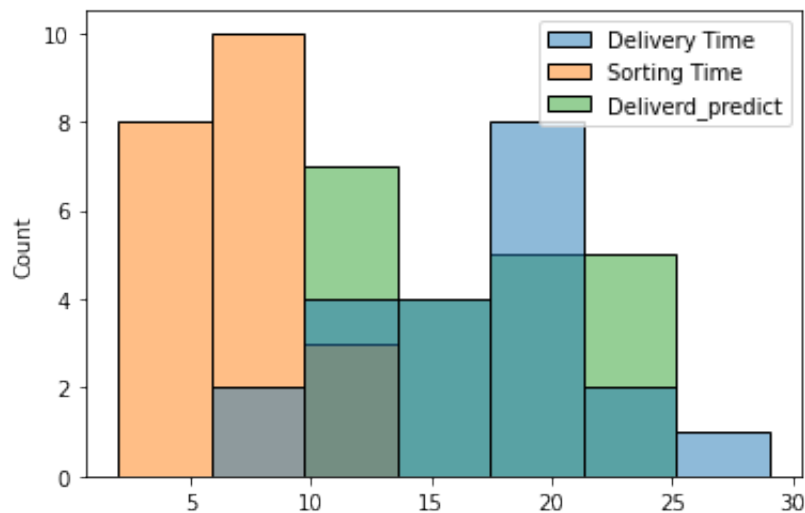
```
In [42]: sns.pairplot(data = df)
```

```
Out[42]: <seaborn.axisgrid.PairGrid at 0x7fe97035c460>
```



```
In [44]: sns.histplot(data = df)
```


```
Out[44]: <AxesSubplot:ylabel='Count'>
```



```
In [5]: sns.distplot(df['Delivery Time'])
```

/opt/anaconda3/lib/python3.9/site-packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)


```
Out[5]: <AxesSubplot:xlabel='Delivery Time', ylabel='Density'>
```



```
In [6]: sns.distplot(df['Sorting Time'])
```

```
/opt/anaconda3/lib/python3.9/site-packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

```
Out[6]: <AxesSubplot:xlabel='Sorting Time', ylabel='Density'>
```



```
In [7]: df.corr()
```

```
Out[7]:
```

	Delivery Time	Sorting Time
Delivery Time	1.000000	0.825997
Sorting Time	0.825997	1.000000

```
In [8]: df['Sorting Time'].describe()
```


```
Out[8]: count    21.000000
mean      6.190476
std       2.542028
min       2.000000
25%       4.000000
50%       6.000000
75%       8.000000
max      10.000000
Name: Sorting Time, dtype: float64
```

```
In [9]: df['Delivery Time'].describe()
```

```
Out[9]: count      21.000000  
mean       16.790952  
std        5.074901  
min         8.000000  
25%        13.500000  
50%        17.830000  
75%        19.750000  
max        29.000000  
Name: Delivery Time, dtype: float64
```

```
In [10]: sns.pairplot(df)
```

```
Out[10]: <seaborn.axisgrid.PairGrid at 0x7fe96f79bdf0>
```



```
In [11]: model = smf.ols("df['Delivery Time']~df['Sorting Time']",data = df)
```

In [12]: `model.summary()`

Out [12]: OLS Regression Results

Dep. Variable:	df['Delivery Time']	R-squared:	0.682
Model:	OLS	Adj. R-squared:	0.666
Method:	Least Squares	F-statistic:	40.80
Date:	Thu, 01 Dec 2022	Prob (F-statistic):	3.98e-06
Time:	02:01:40	Log-Likelihood:	-51.357
No. Observations:	21	AIC:	106.7
Df Residuals:	19	BIC:	108.8
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	6.5827	1.722	3.823	0.001	2.979	10.186
df['Sorting Time']	1.6490	0.258	6.387	0.000	1.109	2.189

Omnibus:	3.649	Durbin-Watson:	1.248
Prob(Omnibus):	0.161	Jarque-Bera (JB):	2.086
Skew:	0.750	Prob(JB):	0.352
Kurtosis:	3.367	Cond. No.	18.3

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


```
In [13]: print(model.summary())
prediction = model.predict(df.iloc[:,1])
```

```

                                OLS Regression Results
=====
Dep. Variable:          df['Delivery Time']    R-squared:
0.682
Model:                                OLS    Adj. R-squared:
0.666
Method:                    Least Squares    F-statistic:
40.80
Date:                    Thu, 01 Dec 2022    Prob (F-statistic):
3.98e-06
Time:                    02:01:40    Log-Likelihood:
-51.357
No. Observations:                21    AIC:
106.7
Df Residuals:                    19    BIC:
108.8
Df Model:                        1
Covariance Type:                nonrobust
=====
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
Intercept                6.5827    1.722    3.823    0.001
2.979    10.186
df['Sorting Time']      1.6490    0.258    6.387    0.000
1.109    2.189
=====
=====
Omnibus:                3.649    Durbin-Watson:
1.248
Prob(Omnibus):          0.161    Jarque-Bera (JB):
2.086
Skew:                  0.750    Prob(JB):
0.352
Kurtosis:              3.367    Cond. No.
18.3
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the error s is correctly specified.

```
In [14]: slr_model=smf.ols("df['Delivery Time']~df['Sorting Time']",data=df)
print(slr_model.summary())
predict=slr_model.predict(df.iloc[:,1])
import matplotlib.pyplot as plt
```

```
plt.scatter(x=df['Sorting Time'],y=df['Delivery Time'])
plt.plot(df['Sorting Time'],predict,color='black')
plt.xlabel("Sorting Time")
plt.ylabel("Delivery Time")
plt.title("slr_model plotting")
plt.show()
```


OLS Regression Results

```
=====
=====
Dep. Variable:          df['Delivery Time']    R-squared:
0.682
Model:                  OLS                  Adj. R-squared:
0.666
Method:                 Least Squares        F-statistic:
40.80
Date:                  Thu, 01 Dec 2022      Prob (F-statistic):
3.98e-06
Time:                  02:01:40             Log-Likelihood:
-51.357
No. Observations:      21                  AIC:
106.7
Df Residuals:          19                  BIC:
108.8
Df Model:              1
Covariance Type:       nonrobust
=====
=====
[0.025    0.975]      coef      std err          t      P>|t|
-----
Intercept              6.5827      1.722      3.823      0.001
2.979    10.186
df['Sorting Time']     1.6490      0.258      6.387      0.000
1.109     2.189
=====
=====
Omnibus:              3.649    Durbin-Watson:
1.248
Prob(Omnibus):        0.161    Jarque-Bera (JB):
2.086
Skew:                 0.750    Prob(JB):
0.352
Kurtosis:             3.367    Cond. No.
18.3
=====
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the error s is correctly specified.

slr_model plotting



```
In [15]: log_model=smf.ols("df['Delivery Time']~np.log(df['Sorting Time'])",
print(log_model.summary())
log_predict=log_model.predict(pd.DataFrame(df['Sorting Time']))

import matplotlib.pyplot as plt
%matplotlib inline
plt.scatter(x=df['Sorting Time'],y=df['Delivery Time'])
plt.plot(df['Sorting Time'],log_predict,color='black')
plt.xlabel("Sorting Time")
plt.ylabel("Delivery Time")
plt.title("log_model plotting")
plt.show()
```

OLS Regression Results

```
=====
=====
Dep. Variable:          df['Delivery Time']    R-squared:
0.695
Model:                                OLS    Adj. R-squared:
0.679
Method:                    Least Squares    F-statistic:
43.39
Date:                    Thu, 01 Dec 2022    Prob (F-statistic):
2.64e-06
Time:                    02:01:41    Log-Likelihood:
-50.912
No. Observations:                21    AIC:
105.8
Df Residuals:                    19    BIC:
107.9
Df Model:                        1
Covariance Type:                nonrobust
=====
=====
```

	coef	std err	t	P
> t	[0.025	0.975]		

```
=====
```

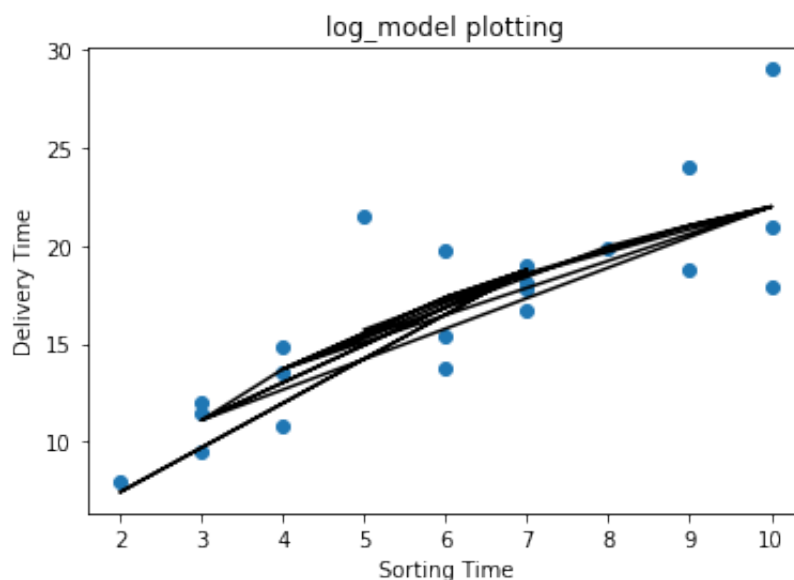
```

-----
Intercept                1.1597        2.455        0.472        0
.642        -3.978        6.297
np.log(df['Sorting Time'])  9.0434        1.373        6.587        0
.000        6.170        11.917
=====
=====
Omnibus:                5.552    Durbin-Watson:
1.427
Prob(Omnibus):          0.062    Jarque-Bera (JB):
3.481
Skew:                   0.946    Prob(JB):
0.175
Kurtosis:              3.628    Cond. No.
9.08
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



```

In [16]: exp_model=smf.ols("np.log(df['Delivery Time'])~(df['Sorting Time'])
print(exp_model.summary())
predict_exp=exp_model.predict(pd.DataFrame(df['Sorting Time']))
pred_exp=np.exp(predict_exp)

import matplotlib.pyplot as plt
%matplotlib inline
plt.scatter(x=df['Sorting Time'],y=df['Delivery Time'])
plt.plot(df['Sorting Time'],np.exp(predict_exp),color='black')
plt.title("Exponential_model plotting")
plt.xlabel("Sorting Time")
plt.ylabel("Delivery Time")
plt.show()

```

OLS Regression Results

```

=====
=====
Dep. Variable:      np.log(df['Delivery Time'])    R-squared:
0.711
Model:              OLS                          Adj. R-squared:
0.696
Method:              Least Squares                F-statistic:
46.73
Date:                Thu, 01 Dec 2022             Prob (F-statistic):
1.59e-06
Time:                02:01:41                     Log-Likelihood:
7.7920
No. Observations:    21                          AIC:
-11.58
Df Residuals:        19                          BIC:
-9.495
Df Model:             1
Covariance Type:     nonrobust
=====
=====

```

		coef	std err	t	P> t
[0.025	0.975]				
Intercept		2.1214	0.103	20.601	0.000
1.906	2.337				
df['Sorting Time']		0.1056	0.015	6.836	0.000
0.073	0.138				


```


=====
=====
Omnibus:              1.238    Durbin-Watson:
1.325
Prob(Omnibus):        0.538    Jarque-Bera (JB):
0.544
Skew:                 0.393    Prob(JB):
0.762
Kurtosis:              3.067    Cond. No.
18.3
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the error s is correctly specified.





```
In [20]: np.sqrt(np.mean((df['Delivery Time']-predict)**2))
```

```
Out[20]: 2.7916503270617654
```

```
In [22]: np.sqrt(np.mean((df['Delivery Time']-predict_exp)**2))
```

```
Out[22]: 14.795516941016686
```

```
In [24]: np.sqrt(np.mean((df['Delivery Time']-log_predict)**2))
```

```
Out[24]: 2.733171476682066
```

```
In [25]: from sklearn.linear_model import LinearRegression
linear_model = LinearRegression()
```

```
In [26]: linear_model=smf.ols("df['Delivery Time']~df['Sorting Time']",data
```

```
In [27]: linear_model=linear_model.fit()
```

```
In [28]: linear_model.params
```

```
Out[28]: Intercept          6.582734
df['Sorting Time']      1.649020
dtype: float64
```

```
In [29]: predicteddata = linear_model.predict(df)
predicteddata
```

```
Out [29]: 0      23.072933
          1      13.178814
          2      16.476853
          3      21.423913
          4      23.072933
          5      16.476853
          6      18.125873
          7      11.529794
          8      23.072933
          9      21.423913
         10      19.774893
         11      13.178814
         12      18.125873
         13      11.529794
         14      11.529794
         15      13.178814
         16      16.476853
         17      18.125873
         18       9.880774
         19      18.125873
         20      14.827833
dtype: float64
```

```
In [30]: print(model.tvalues, '\n', model.pvalues)
```

```
Intercept          3.823349
df['Sorting Time']  6.387447
dtype: float64
Intercept          0.001147
df['Sorting Time']  0.000004
dtype: float64
```

```
In [31]: #R squared values
(model.rsquared,model.rsquared_adj)
```

```
Out [31]: (0.6822714748417231, 0.6655489208860244)
```

In [32]: `model.summary()`

Out [32]: OLS Regression Results

Dep. Variable:	df['Delivery Time']	R-squared:	0.682
Model:	OLS	Adj. R-squared:	0.666
Method:	Least Squares	F-statistic:	40.80
Date:	Thu, 01 Dec 2022	Prob (F-statistic):	3.98e-06
Time:	02:04:15	Log-Likelihood:	-51.357
No. Observations:	21	AIC:	106.7
Df Residuals:	19	BIC:	108.8
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	6.5827	1.722	3.823	0.001	2.979	10.186
df['Sorting Time']	1.6490	0.258	6.387	0.000	1.109	2.189

Omnibus:	3.649	Durbin-Watson:	1.248
Prob(Omnibus):	0.161	Jarque-Bera (JB):	2.086
Skew:	0.750	Prob(JB):	0.352
Kurtosis:	3.367	Cond. No.	18.3

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [33]: `Deliverd = model.predict()`
`Deliverd`

Out [33]: array([23.07293294, 13.17881356, 16.47685335, 21.42391304, 23.07293294, 16.47685335, 18.12587325, 11.52979366, 23.07293294, 21.42391304, 19.77489315, 13.17881356, 18.12587325, 11.52979366, 11.52979366, 13.17881356, 16.47685335, 18.12587325, 9.88077377, 18.12587325, 14.82783346])


```
In [34]: Deliverd_predict=pd.DataFrame(Deliverd, columns=['Salary_hike'])  
Deliverd_predict
```

Out [34]:

	Salary_hike
0	23.072933
1	13.178814
2	16.476853
3	21.423913
4	23.072933
5	16.476853
6	18.125873
7	11.529794
8	23.072933
9	21.423913
10	19.774893
11	13.178814
12	18.125873
13	11.529794
14	11.529794
15	13.178814
16	16.476853
17	18.125873
18	9.880774
19	18.125873
20	14.827833

```
In [35]: Deliverd_predict=pd.DataFrame(Deliverd, columns=['Salary_hike'])  
Deliverd_predict
```

Out [35]:

	Salary_hike
0	23.072933
1	13.178814
2	16.476853
3	21.423913
4	23.072933
5	16.476853
6	18.125873
7	11.529794
8	23.072933
9	21.423913
10	19.774893
11	13.178814
12	18.125873
13	11.529794
14	11.529794
15	13.178814
16	16.476853
17	18.125873
18	9.880774
19	18.125873
20	14.827833

```
In [36]: df['Deliverd_predict'] = Deliverd_predict  
df
```

Out[36]:

	Delivery Time	Sorting Time	Deliverd_predict
0	21.00	10	23.072933
1	13.50	4	13.178814
2	19.75	6	16.476853
3	24.00	9	21.423913
4	29.00	10	23.072933
5	15.35	6	16.476853
6	19.00	7	18.125873
7	9.50	3	11.529794
8	17.90	10	23.072933
9	18.75	9	21.423913
10	19.83	8	19.774893
11	10.75	4	13.178814
12	16.68	7	18.125873
13	11.50	3	11.529794
14	12.03	3	11.529794
15	14.88	4	13.178814
16	13.75	6	16.476853
17	18.11	7	18.125873
18	8.00	2	9.880774
19	17.83	7	18.125873
20	21.50	5	14.827833

In []:

In []: