

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import statsmodels.formula.api as smf
import matplotlib.pyplot as plt
from sklearn import linear_model
```

```
In [2]: df = pd.read_csv("Salary_Data.csv")
df
```

Out[2]:

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0
5	2.9	56642.0
6	3.0	60150.0
7	3.2	54445.0
8	3.2	64445.0
9	3.7	57189.0
10	3.9	63218.0
11	4.0	55794.0
12	4.0	56957.0
13	4.1	57081.0
14	4.5	61111.0
15	4.9	67938.0
16	5.1	66029.0
17	5.3	83088.0
18	5.9	81363.0
19	6.0	93940.0
20	6.8	91738.0
21	7.1	98273.0
22	7.9	101302.0
23	8.2	113812.0
24	8.7	109431.0

25	9.0	105582.0
26	9.5	116969.0
27	9.6	112635.0
28	10.3	122391.0
29	10.5	121872.0

```
In [3]: df.head()
```

```
Out[3]:
```

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0

```
In [4]: df.dtypes
```

```
Out[4]: YearsExperience    float64
Salary                  float64
dtype: object
```

```
In [6]: df.describe()
```

```
Out[6]:
```

	YearsExperience	Salary
count	30.000000	30.000000
mean	5.313333	76003.000000
std	2.837888	27414.429785
min	1.100000	37731.000000
25%	3.200000	56720.750000
50%	4.700000	65237.000000
75%	7.700000	100544.750000
max	10.500000	122391.000000

```
In [7]: df.isnull().sum()
```

```
Out[7]: YearsExperience    0
Salary                  0
dtype: int64
```


In [8]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   YearsExperience  30 non-null    float64
1   Salary           30 non-null    float64
dtypes: float64(2)
memory usage: 608.0 bytes
```

In [9]: `sns.distplot(df['Salary'])`

```
/opt/anaconda3/lib/python3.9/site-packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```


Out [9]: `<AxesSubplot:xlabel='Salary', ylabel='Density'>`



```
In [10]: sns.distplot(df['YearsExperience'])
```

```
/opt/anaconda3/lib/python3.9/site-packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

```
Out[10]: <AxesSubplot:xlabel='YearsExperience', ylabel='Density'>
```



```
In [11]: df.corr()
```

```
Out[11]:
```

	YearsExperience	Salary
YearsExperience	1.000000	0.978242
Salary	0.978242	1.000000

```
In [12]: df=df.rename({'YearsExperience':'years_exp', 'Salary':'salary_hike'})
df
```


```
Out[12]:
```

	years_exp	salary_hike
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0
5	2.9	56642.0
6	3.0	60150.0
7	3.2	54445.0

8	3.2	64445.0
9	3.7	57189.0
10	3.9	63218.0
11	4.0	55794.0
12	4.0	56957.0
13	4.1	57081.0
14	4.5	61111.0
15	4.9	67938.0
16	5.1	66029.0
17	5.3	83088.0
18	5.9	81363.0
19	6.0	93940.0
20	6.8	91738.0
21	7.1	98273.0
22	7.9	101302.0
23	8.2	113812.0
24	8.7	109431.0
25	9.0	105582.0
26	9.5	116969.0
27	9.6	112635.0
28	10.3	122391.0
29	10.5	121872.0

```
In [13]: plt.scatter(df.years_exp,df.salary_hike)
```

```
Out[13]: <matplotlib.collections.PathCollection at 0x7fda462d5e20>
```



```
In [14]: slr_model=smf.ols("df['salary_hike']~df['years_exp']",data=df).fit(
print(slr_model.summary())
predict=slr_model.predict(df.iloc[:,0])

plt.scatter(df['years_exp'],df['salary_hike'])
plt.plot(df['years_exp'],predict)
plt.title("SLR Data set")
plt.xlabel("Experience in Years")
plt.ylabel("Salary")
plt.show()
```

OLS Regression Results

```
=====
=====
Dep. Variable:          df['salary_hike']    R-squared:
0.957
Model:                                OLS    Adj. R-squared:
0.955
Method:                    Least Squares    F-statistic:
622.5
Date:                    Thu, 01 Dec 2022    Prob (F-statistic):
1.14e-20
Time:                    02:19:05           Log-Likelihood:
-301.44
No. Observations:                30         AIC:
606.9
Df Residuals:                    28         BIC:
609.7
Df Model:                        1
Covariance Type:                nonrobust
=====
=====
```

	coef	std err	t	P> t	[
--	------	---------	---	------	---

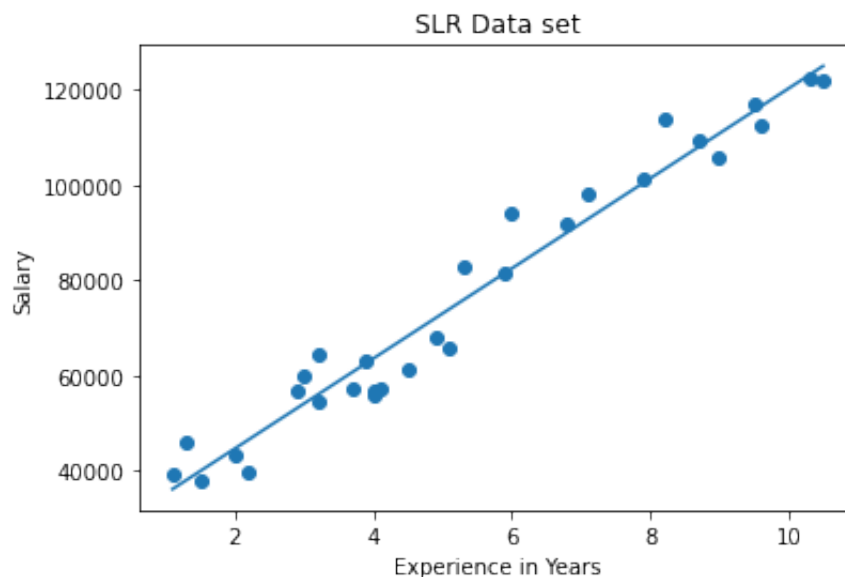
0.025 0.975]

```
-----
-----
Intercept      2.579e+04   2273.053   11.347   0.000   2.1
1e+04   3.04e+04
df['years_exp'] 9449.9623   378.755   24.950   0.000   867
4.119   1.02e+04
=====
=====
```

```
Omnibus:                2.140   Durbin-Watson:
1.648
Prob(Omnibus):          0.343   Jarque-Bera (JB):
1.569
Skew:                   0.363   Prob(JB):
0.456
Kurtosis:               2.147   Cond. No.
13.2
=====
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the error s is correctly specified.



```
In [15]: log_model=smf.ols("df['salary_hike']~np.log(df['years_exp'])", data
print(log_model.summary())
log_predict=log_model.predict(pd.DataFrame(df['years_exp']))

plt.scatter(df['years_exp'],df['salary_hike'])
plt.plot(df['years_exp'],log_predict)
plt.title("log Data model")
plt.xlabel("Experience in Years")
plt.ylabel("Salary")
plt.show()
```

OLS Regression Results

```

=====
=====
Dep. Variable:          df['salary_hike']    R-squared:
0.854
Model:                                OLS    Adj. R-squared:
0.849
Method:                    Least Squares    F-statistic:
163.6
Date:                    Thu, 01 Dec 2022    Prob (F-statistic):
3.25e-13
Time:                    02:19:18           Log-Likelihood:
-319.77
No. Observations:                30    AIC:
643.5
Df Residuals:                    28    BIC:
646.3
Df Model:                        1
Covariance Type:                nonrobust
=====
=====


```


		coef	std err	t	P> t

Intercept		1.493e+04	5156.226	2.895	0.00
7	4365.921	2.55e+04			
np.log(df['years_exp'])		4.058e+04	3172.453	12.792	0.00
0	3.41e+04	4.71e+04			
=====					
=====					
Omnibus:		1.094		Durbin-Watson:	
0.512					
Prob(Omnibus):		0.579		Jarque-Bera (JB):	
0.908					
Skew:		0.156		Prob(JB):	
0.635					
Kurtosis:		2.207		Cond. No.	
5.76					
=====					
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the error s is correctly specified.





```
In [16]: exp_model=smf.ols("np.log(df['salary_hike'])~df['years_exp']", data
print(exp_model.summary())
exp_predict=exp_model.predict(pd.DataFrame(df['years_exp']))
pred_exp=np.exp(exp_predict)

plt.scatter(df['years_exp'],df['salary_hike'])
plt.plot(df['years_exp'],np.exp(exp_predict))
plt.title("Exponential_model plotting")
plt.xlabel("Experience in Years")
plt.ylabel("Salary")
plt.show()
```

OLS Regression Results

```
=====
=====
Dep. Variable:      np.log(df['salary_hike'])    R-squared:
0.932
Model:              OLS                        Adj. R-squared:
0.930
Method:            Least Squares              F-statistic:
383.6
Date:              Thu, 01 Dec 2022            Prob (F-statistic):
7.03e-18
Time:              02:19:28                    Log-Likelihood:
28.183
No. Observations:  30                        AIC:
-52.37
Df Residuals:      28                        BIC:
-49.56
Df Model:          1
Covariance Type:   nonrobust
=====
=====

```

	coef	std err	t	P> t	[
Intercept	10.5074	0.038	273.327	0.000	1
df['years_exp']	0.1255	0.006	19.585	0.000	

```
=====
=====
Omnibus:          0.826    Durbin-Watson:
1.438
```

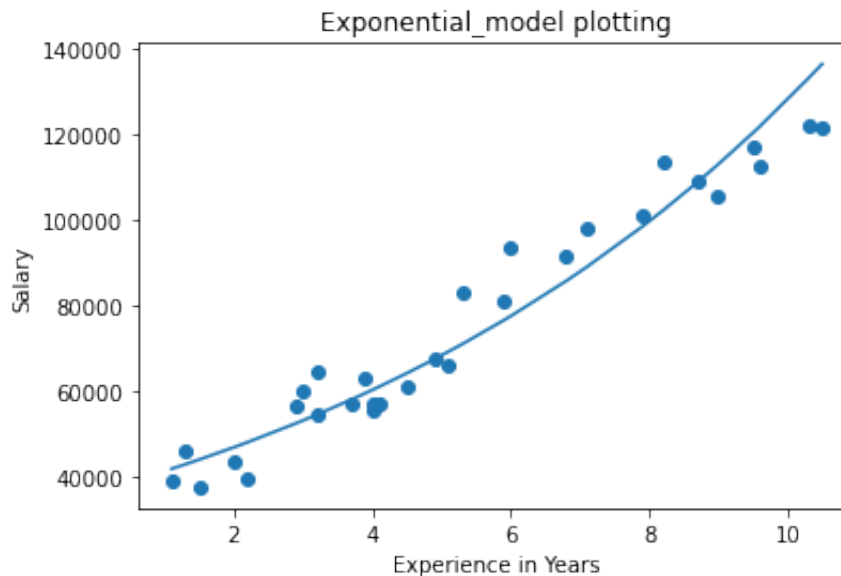
Prob(Omnibus):	0.661	Jarque-Bera (JB):
0.812		
Skew:	0.187	Prob(JB):
0.666		
Kurtosis:	2.286	Cond. No.
13.2		

=====

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the error s is correctly specified.



```
In [17]: slr_regression=linear_model.LinearRegression()
slr_regression.fit(df[['years_exp']],df.salary_hike)

flag=True
while(flag):
    myinput=float(input("Enter Years of Experience"))
    myoutput=slr_regression.predict([[myinput]])
    print("Salary Hike precited using SLR is: ",myoutput)
    flag=int(input("press 1 to continue or press 0 to exit: "))
```

```
Enter Years of Experience5
Salary Hike precited using SLR is: [73042.01180594]
press 1 to continue or press 0 to exit: 0
```

```
In [18]: np.sqrt(np.mean((df.years_exp-predict)**2))
```

```
Out[18]: 80440.84508275457
```

```
In [19]: np.sqrt(np.mean((df.years_exp-predict)**2))
```

```
Out[19]: 80440.84508275457
```

```
In [20]: np.sqrt(np.mean((df.years_exp-log_predict)**2))
```

```
Out[20]: 79974.15496099806
```

```
In [ ]:
```